

The Investment Philosophy

of the AI Hedge Fund

Volume I - Version 1.0

Prepared as the founding investment doctrine for an AI-native research and portfolio system focused on AI, semiconductors, digital infrastructure, and disciplined capital allocation.

Status: Working draft for research, paper trading, and process design. This is not personalized financial advice, legal advice, tax advice, or a recommendation to buy or sell any security.

Date: June 29, 2026

Table of Contents

Preface

Chapter 1 - What Investing Really Is

Chapter 2 - The Purpose of Capital

Chapter 3 - Our Definition of Risk

Chapter 4 - The Search for Edge

Chapter 5 - Truth Above Conviction

Chapter 6 - Evidence Before Opinion

Chapter 7 - AI as an Investment Partner

Chapter 8 - Building a Research Culture

Chapter 9 - Circle of Competence

Chapter 10 - Portfolio Philosophy

Chapter 11 - When Not to Invest

Chapter 12 - Continuous Improvement

Appendices

Bibliography

Preface

This document exists to answer a simple but brutal question: what does this firm actually believe?

A trading system without a philosophy is just a machine for producing activity. It can generate charts, alerts, backtests, confidence scores, and smart-sounding recommendations, but it has no internal standard for deciding whether any of that activity is intelligent. The purpose of this document is to create that standard.

The AI Hedge Fund is not imagined as a black-box stock-picking toy. It is designed as an AI-native investment research firm: a disciplined system of specialist agents, structured evidence, independent debate, risk controls, and performance measurement. Its objective is not to predict every market move. Its objective is to make better decisions over time.

This philosophy draws from several durable traditions: Graham-style margin of safety, Buffett-style business quality, Marks-style risk awareness and second-level thinking, Dalio-style macro awareness and principles, Markowitz-style portfolio construction, Fama-French factor research, momentum evidence from Jegadeesh, Titman, and Carhart, behavioral finance from Kahneman and Thaler, and quantitative validation associated with systematic investors such as Renaissance Technologies. The goal is not to imitate any one school. The goal is to synthesize what survives contact with reality.

The firm's first domain is artificial intelligence, semiconductors, cloud infrastructure, data centers, and the picks-and-shovels businesses that enable modern compute. That domain is volatile, crowded, and prone to hype. It is also one of the clearest areas where technology, capital spending, supply chains, and productivity may reshape the economy. That combination makes it attractive, but dangerous. The philosophy must therefore be both ambitious and skeptical.

We will not confuse intelligence with complexity. We will not confuse data with truth. We will not confuse confidence with accuracy. We will not confuse a rising stock price with a good decision. And we will not let the system trade simply because it is capable of trading.

The highest expression of the system is not a flashy buy signal. It is a clear, evidence-backed recommendation that includes reasons to act, reasons not to act, defined risk, confidence, and an honest statement of what is unknown.

Chapter 1 - What Investing Really Is

Investing is the act of allocating capital today in expectation of a future stream of cash flows, value creation, or risk-adjusted return. Speculation is the act of buying something because someone else may pay more for it later. The two can overlap in practice, but they are not the same mental activity.

The AI Hedge Fund must be able to trade, but it must think like an investor first. That distinction matters because a system that only asks, “Can the price go up today?” becomes dependent on noise. A system that asks, “What is the asset worth, why might the market be wrong, and what evidence would prove us wrong?” can survive multiple market regimes.

Markets are information-processing systems. They absorb facts, expectations, emotion, positioning, liquidity, forced selling, regulatory changes, interest rates, geopolitics, and human error. Sometimes markets are extraordinarily efficient. Sometimes they are absurd. The firm’s philosophy begins with humility: the market is not always right, but it is never obligated to agree with us on our timeline.

A business becomes valuable when it can generate cash flows above the cost of capital for a long time. A stock becomes a good investment only when the price paid leaves room for attractive return after considering risk, uncertainty, and alternatives. These are different ideas. Great companies can be terrible investments at the wrong price. Mediocre companies can be attractive trades at extreme pessimism. The system must understand both business quality and price.

In AI and semiconductors, the temptation is to treat every important company as an automatic buy. That is lazy. Importance is not the same as undervaluation. NVIDIA, TSMC, ASML, Broadcom, AMD, Arm, Micron, and the broader semiconductor supply chain may be central to the future economy, but the future can still be over-discounted, supply can disappoint, margins can compress, export controls can bite, and capital cycles can reverse. The system’s job is to separate secular importance from investable opportunity.

Principles

- We invest in expected value, not stories.
- We separate company quality from stock attractiveness.
- We assume the market may be efficient enough to punish laziness and inefficient enough to reward disciplined research.
- We prefer repeatable processes over impressive one-off predictions.

Agent Implications

Every agent must label its output as one of three types: fact, interpretation, or recommendation. A fact is observable. An interpretation is a judgment about meaning. A recommendation is a proposed action. Mixing those categories is how a research process turns into a sales pitch.

Chapter 2 - The Purpose of Capital

Capital is stored optionality. It can be deployed into productive businesses, preserved for future opportunities, used to absorb volatility, or wasted on low-quality action. The firm's purpose is to compound capital intelligently, not to maximize excitement.

Compounding requires two ingredients: return and survival. Return without survival is a casino run. Survival without return is stagnation. The operating philosophy must pursue both, with survival taking priority because capital that is destroyed cannot compound.

Long-term investing is not the same as holding forever. It is the willingness to let a valid thesis mature while constantly checking whether the thesis remains valid. The system should not sell just because a stock is down, nor should it hold just because it originally liked the company. The reason for owning an asset must be alive. If the reason dies, the position must be reconsidered.

The AI Hedge Fund should also distinguish between core capital and tactical capital. Core capital is allocated to durable businesses, broad themes, and long-term compounding. Tactical capital is allocated to shorter-term opportunities where evidence suggests a favorable setup. The two should not be managed with the same rules. A long-term AI infrastructure thesis may tolerate volatility. A one-day trade cannot.

Capital Buckets

Bucket	Purpose	Typical Horizon	Primary Risk Control
Core Compounders	Own durable businesses and themes with long runways	Years	Valuation discipline and thesis monitoring
Swing Opportunities	Exploit medium-term mispricing, earnings revisions, or momentum	Days to months	Stops, thesis invalidation, and position limits
Intraday / Tactical	Act on high-probability setups with defined levels	Minutes to days	Predefined loss limits and no overnight assumption
Cash / Reserves	Preserve optionality and reduce forced selling	Always available	Opportunity-cost review

Principles

- Cash is not laziness; cash is ammunition.
- The system must know whether a recommendation belongs to core, swing, or tactical capital.
- Capital allocation decisions must consider opportunity cost.
- The firm compounds by avoiding stupid losses as much as by finding brilliant wins.

Chapter 3 - Our Definition of Risk

Risk is not simply volatility. Volatility is movement. Risk is the probability and magnitude of unacceptable loss, including permanent impairment of capital, forced liquidation, strategy failure, and behavioral error.

Markowitz gave finance a powerful language for thinking about portfolio risk: expected return, variance, covariance, and diversification. That framework remains foundational because it forces investors to evaluate assets in relation to the whole portfolio rather than as isolated bets. But variance is not enough. A portfolio can appear diversified across tickers while being concentrated in the same underlying factor: AI hype, long-duration growth, falling interest rates, semiconductor capex, or U.S. mega-cap technology.

For this firm, risk has multiple layers: security-level risk, thesis risk, valuation risk, factor risk, liquidity risk, concentration risk, macro risk, regulatory risk, model risk, execution risk, and behavioral risk. A proper risk manager looks across all of them.

The AI Hedge Fund will not treat a high-confidence idea as permission to oversize a position. Confidence is an input to sizing, not a substitute for risk control. The market does not care how confident the agents are. A wrong high-confidence trade is still wrong.

Risk Taxonomy

Risk Type	Definition	Required Control
Thesis Risk	The core reason for owning/trading the asset is wrong	Explicit thesis and invalidation criteria
Valuation Risk	The company performs well but the stock was too expensive	Expected return analysis and margin of safety
Factor Risk	Multiple holdings depend on the same driver	Factor exposure monitoring
Liquidity Risk	Position cannot be exited cleanly	Volume, spread, and order-size limits
Model Risk	The agent or backtest is wrong, stale, or overfit	Out-of-sample testing and human review
Behavioral Risk	FOMO, revenge trading, anchoring, confirmation bias	Devil's Advocate and no-trade principle

The Risk Manager's Authority

The Risk Manager has veto authority. This is not symbolic. If the Risk Manager rejects a trade due to position size, correlation, drawdown, event risk, poor liquidity, or insufficient evidence, the trade does not occur. In a human firm, this rule prevents ego. In an AI-native firm, it prevents the system from optimizing for activity instead of survival.

Principles

- Risk is what can prevent the strategy from continuing.
- Volatility is tolerated only when compensated by expected return.
- Every trade must define maximum acceptable loss before entry.
- The system must never average down without a predefined thesis and risk budget.
- No agent may override risk controls by appealing to conviction.

Chapter 4 - The Search for Edge

An edge is a repeatable reason to expect superior risk-adjusted returns after costs, taxes, slippage, and mistakes. A story is not an edge. A chart pattern is not automatically an edge. A popular theme is not an edge. A model output is not an edge until it has been validated.

The firm recognizes four major categories of edge: information edge, analytical edge, behavioral edge, and execution edge. Most retail investors have weak information edge because public information is widely available. The firm should therefore prioritize analytical, behavioral, and process edge: synthesizing public information faster and more objectively, avoiding common emotional errors, and logging outcomes with discipline.

In AI and semiconductors, information often arrives in fragments: hyperscaler capex guidance, foundry utilization, HBM supply, advanced packaging bottlenecks, export controls, earnings-call language, channel checks, power constraints, and supply-chain announcements. The firm's edge may come from connecting those fragments before the market fully appreciates their combined implication.

Edges decay. Once a signal becomes popular, arbitrated, or regime-dependent, its usefulness can fall. The system must therefore treat every strategy as provisional. Backtesting is not proof. It is evidence. Live and paper-trading performance provide additional evidence, but even successful performance must be evaluated for luck, hidden factor exposure, and changing market structure.

Edge Types

Edge	Example	Primary Test
Information Edge	Faster access to filings, transcripts, economic releases, or specialized datasets	Timeliness, uniqueness, reliability
Analytical Edge	Better synthesis of fundamentals, macro, positioning, and technicals	Forward outcome tracking
Behavioral Edge	Avoiding FOMO, panic, anchoring, and consensus chasing	Trade journal review
Execution Edge	Better order placement, sizing, and discipline	Slippage and fill analysis
Technological Edge	Automated monitoring, agent collaboration, structured memory	Process speed and accuracy

Principles

- Every strategy must name its expected edge.
- An edge must survive costs and realistic execution.

- Popular signals deserve extra skepticism.
- When the system cannot identify a clear edge, the correct recommendation is no trade.

Chapter 5 - Truth Above Conviction

Conviction is useful only when it is earned. Unearned conviction is how investors turn uncertainty into losses. The AI Hedge Fund will prefer calibrated uncertainty over false precision.

The firm's reasoning standard is Bayesian in spirit: start with a prior belief, update that belief as evidence arrives, and change the conclusion when the weight of evidence changes. A thesis is not a marriage. It is a probability estimate. If the facts change, the probability changes.

Agents must not be rewarded for sounding confident. They must be rewarded for being accurate, calibrated, and honest. An agent that says "insufficient information" when evidence is weak is doing its job. An agent that fabricates certainty to satisfy the user is a liability.

This is especially important in fast-moving AI markets. Narratives can run far ahead of realized cash flows. Management teams can use the language of AI without meaningful revenue. A chip shortage can turn into oversupply. A capex boom can become a digestion period. The firm must update quickly, but not chase every headline.

Confidence Framework

Confidence	Meaning	Action Standard
0-40%	Evidence is weak, conflicting, or incomplete	No trade; research only
41-60%	Possible edge but uncertainty is high	Watchlist or very small sizing only
61-75%	Evidence supports action but risks remain meaningful	Eligible for controlled position
76-90%	Multiple independent signals align	Eligible for higher conviction sizing within limits
91-100%	Rare; overwhelming evidence and favorable risk/reward	Still subject to Risk Manager veto

Principles

- A recommendation without confidence is incomplete.
- Confidence must be calibrated against historical outcomes.
- The system must update beliefs when evidence changes.
- The words "always," "guaranteed," and "can't lose" are banned from investment reasoning.

Chapter 6 - Evidence Before Opinion

The firm's evidence hierarchy begins with primary sources: filings, earnings releases, transcripts, audited financials, official economic data, exchange data, and broker-confirmed execution records. Secondary sources such as news articles, analyst notes, and expert commentary can be useful, but they must not be treated as primary evidence unless independently verified.

For U.S. company filings and XBRL data, the SEC provides EDGAR application programming interfaces, including company submissions and extracted company facts. This matters because a serious system should not rely only on scraped summaries when primary filings are available. [16]

Good data is not the same as more data. Bad data at scale is just industrialized confusion. Each dataset must be scored for source quality, latency, completeness, survivorship bias, revision history, licensing, and relevance. A cheap data source can be excellent for research but unacceptable for execution. A premium data source can still be useless if it does not improve decisions.

The firm must also distinguish evidence from evidence-shaped content. Social media chatter, rumors, influencer posts, and unsourced "smart money is buying" claims are not evidence unless tied to verifiable data such as filings, order-flow data, or reliable market records.

Evidence Hierarchy

Tier	Evidence Type	Weight
Tier 1	Primary filings, audited statements, exchange data, official economic releases	Highest
Tier 2	Earnings transcripts, investor presentations, company press releases	High but subject to management bias
Tier 3	Reputable financial news, analyst research, industry reports	Useful but must be verified
Tier 4	Alternative data, options flow, dark-pool indicators, social sentiment	Potentially useful but prone to noise
Tier 5	Unverified claims, rumors, anonymous posts	Normally zero weight

Principles

- Every important claim must be traceable.
- Primary data outranks commentary.
- Data quality is part of investment quality.
- If the source is unclear, the confidence score must fall.

- Agents must say what data is missing.

Chapter 7 - AI as an Investment Partner

AI is an accelerator, not an oracle. It can read more documents than a human, monitor more signals, summarize faster, and detect patterns across messy information. It can also hallucinate, overfit, anchor on flawed prompts, miss causal structure, and produce persuasive nonsense. The philosophy must embrace both truths.

The correct role for AI is not “make me rich.” The correct role is: gather evidence, structure arguments, challenge assumptions, estimate probabilities, monitor changes, and improve the research process. Human oversight remains essential, especially before live execution.

Modern agent frameworks make it practical to build applications where agents use instructions, tools, handoffs, guardrails, and structured outputs. OpenAI’s official documentation describes agents as applications that plan, call tools, collaborate across specialists, and maintain enough state to complete multi-step work. [17] That is exactly the architecture this firm should use: specialist agents with narrow mandates, shared memory, explicit controls, and a supervising CIO agent.

Autonomous trading should be the last step, not the first. Broker APIs such as Alpaca and Interactive Brokers can support automated or semi-automated workflows, and paper trading environments allow simulated testing before real capital is at risk. Alpaca’s own paper-trading documentation emphasizes simulated activity rather than live financial risk, while Interactive Brokers documents API access for trading, account management, and portfolio information. [18] [19] The firm should use these capabilities cautiously: recommendations first, paper trading second, small live trading only after evidence justifies it.

AI Role Definition

AI Should Do	AI Should Not Do
Read filings, transcripts, and data feeds	Invent facts or cite nonexistent sources
Generate structured daily research	Trade live capital without controls
Compare agent disagreement	Optimize for activity
Log outcomes and calibration	Hide uncertainty
Identify missing information	Override risk policy

Principles

- AI augments judgment; it does not replace accountability.
- Tool use must be auditable.
- Every autonomous action requires a permission boundary.
- The system should be designed for graceful failure, not perfect prediction.

Chapter 8 - Building a Research Culture

Culture sounds soft until money is lost. Then culture becomes the difference between learning and denial. The AI Hedge Fund needs a research culture built around intellectual honesty, structured disagreement, and post-decision review.

The Devil's Advocate is not a gimmick. It is a required function. Whenever consensus forms, the Devil's Advocate must attack the thesis, search for base-rate failures, identify hidden assumptions, and explain why the recommendation may be wrong. Consensus without challenge is not discipline; it is groupthink with better formatting.

The Research Director has a different role: improve the system itself. It asks whether agents are adding value, which datasets are missing, which prompts are producing weak analysis, which recommendations were poorly calibrated, and what experiments should be run next.

The firm should celebrate killed ideas. A rejected bad trade is a win. A no-trade decision that avoids a low-quality setup is a win. An agent that identifies missing evidence is not being negative; it is protecting the process.

Research Norms

- Attack ideas, not agents.
- Prefer specific objections to vague caution.
- Record disagreement rather than smoothing it away.
- Reward agents for accurate uncertainty.
- Treat mistakes as data, not embarrassment.

Investment Committee Standard

The final recommendation should include: market regime, opportunity, thesis, evidence, entry plan, exit plan, risk, confidence, dissenting views, missing information, and no-trade alternative. If the system cannot produce this cleanly, it is not ready to trade.

Chapter 9 - Circle of Competence

A circle of competence is not a list of things the firm likes. It is the boundary between what the firm can analyze with discipline and what it merely finds interesting.

The initial circle of competence is AI and semiconductor infrastructure. This includes AI accelerators, foundries, semiconductor equipment, advanced packaging, high-bandwidth memory, networking, cloud capex, data centers, power infrastructure, and software platforms that monetize AI adoption. The circle may expand, but only after the firm can define the key value drivers, data sources, risks, and historical cycles of the new domain.

The danger of a thematic fund is narrative creep. Everything eventually claims to be AI. The system must ask whether AI is central to the economics of the business or merely attached to the investor relations deck. If the revenue, margins, capex, or competitive position do not meaningfully depend on AI, the company may be outside the core thesis.

Initial Coverage Universe

Domain	Examples of What to Study
Compute	GPUs, ASICs, accelerators, CPUs, inference chips
Manufacturing	Foundries, lithography, deposition, etch, inspection
Memory and Packaging	HBM, DRAM, NAND, CoWoS-like advanced packaging, substrates
Networking	Switching, optical interconnect, data-center networking
Cloud and Data Centers	Hyperscaler capex, utilization, power, cooling, real estate
Software and Platforms	AI monetization, developer ecosystems, enterprise adoption

Expansion Criteria

- The firm can name the industry's key value drivers.
- The firm can identify primary data sources.
- The firm can explain the cycle history and failure modes.
- The firm can build a benchmark universe.
- The firm can evaluate whether the new domain improves portfolio risk-adjusted return.

Chapter 10 - Portfolio Philosophy

A portfolio is not a pile of good ideas. It is a system of exposures. The same portfolio can look diversified by ticker and dangerously concentrated by economic driver. This is especially true in AI and semiconductors, where many companies rise and fall with the same assumptions about capex, rates, growth, liquidity, and risk appetite.

The portfolio philosophy combines three ideas. First, from modern portfolio theory, evaluate assets by their contribution to total portfolio risk, not in isolation. Second, from value and quality investing, prefer businesses with durable economics and a margin of safety. Third, from quantitative investing, measure exposures and validate whether chosen signals actually improve outcomes.

The firm should maintain benchmarks. For AI and semiconductors, relevant comparisons might include QQQ, SMH, SOXX, broad market indexes, and an equal-weight custom universe. The system must not declare victory because a trade made money during a bull market. It must ask whether it added value relative to a simple benchmark after costs and risk.

Position Sizing

Position sizing should reflect expected edge, confidence, volatility, liquidity, correlation, event risk, and current drawdown. A high-upside idea with poor liquidity may deserve a smaller position than a lower-upside idea with cleaner risk. The firm should size positions so that a single mistake cannot derail the strategy.

Portfolio Rules

- No position exists outside a defined bucket: core, swing, tactical, hedge, or cash.
- Every position must have an owner agent and a thesis file.
- Correlated exposure must be measured at the portfolio level.
- Benchmark-relative performance must be reported.
- Tax and transaction costs must be included when evaluating active trading.

Chapter 11 - When Not to Invest

The no-trade principle is one of the firm's most important beliefs. A system that feels pressure to produce a trade every day will eventually manufacture reasons. Professional discipline often means waiting while nothing happens.

No-trade decisions should be explicit, not passive. "No trade" means the system evaluated the available evidence and concluded that expected edge did not justify risk. That decision should be logged like any other recommendation because it may later prove valuable. Avoided losses are part of performance.

The system should avoid trading when evidence is weak, key data is missing, spreads are poor, event risk is binary, position correlation is already high, the user is emotionally biased, or the market is behaving outside tested regimes. The ability to say no is an edge because many participants cannot.

Automatic No-Trade Triggers

- The Risk Manager vetoes the trade.
- The recommendation lacks a clear invalidation point.
- Confidence is below the action threshold.
- The thesis depends on unverified rumor.
- The position would breach concentration or drawdown limits.
- Major scheduled event risk is present and not explicitly part of the strategy.
- The trade exists mainly because the system feels compelled to produce activity.

Principle

The firm would rather miss a good trade than take a bad one for a stupid reason.

Chapter 12 - Continuous Improvement

The AI Hedge Fund must learn, but it must not self-modify recklessly. Improvement should follow a research process: identify weakness, propose change, test historically where possible, paper trade where necessary, compare against baseline, and adopt only if evidence improves the system.

Every recommendation should become a data point. The log should include date, market regime, ticker, thesis, agent votes, confidence, entry, stop, target, outcome after multiple horizons, benchmark return, and notes on whether the decision was good regardless of outcome. A profitable bad decision is still bad process. An unprofitable good decision may still be acceptable if the expected value was sound.

Agent performance must be measured by calibration and contribution. Did the agent improve the final decision? Did its high-confidence calls work better than its low-confidence calls? Did it repeatedly miss the same risk? Did it produce useful dissent? Did it rely on stale or weak data?

Review Cadence

Cadence	Review Questions
Daily	Did the system follow its process? Were any recommendations unclear?
Weekly	Which agents added value? Which signals failed? What data was missing?
Monthly	Did the strategy outperform benchmarks after costs? Were confidence scores calibrated?
Quarterly	Should any agent, data source, or strategy be upgraded, retired, or reweighted?
Annually	Does the philosophy still hold? Has the market structure changed?

Principles

- The market is the final judge, but process quality matters.
- The system improves through measured feedback, not vibes.
- No prompt, model, or dataset is sacred.
- Changes require evidence, not excitement.

Appendices

Appendix A - Decision Checklist

1. What is the opportunity?
2. What is the thesis?
3. What are the facts supporting it?
4. What is interpretation rather than fact?
5. What would prove the thesis wrong?
6. What is the entry, stop, target, and time horizon?
7. What is the expected edge?
8. What is the portfolio impact?
9. What does the Devil's Advocate say?
10. Does the Risk Manager approve?
11. Is no trade better?

Appendix B - Research Checklist

12. Review primary filings and recent earnings materials.
13. Compare current thesis to historical financial performance.
14. Analyze valuation relative to growth, margins, and peers.
15. Check macro backdrop: rates, dollar, liquidity, economic data.
16. Review sector-specific drivers: capex, supply, demand, regulation.
17. Review technical structure and liquidity.
18. Review options flow and positioning only as supporting evidence.
19. Identify missing data and unresolved risks.
20. Score confidence and calibration.
21. Prepare recommendation in structured format.

Appendix C - Agent Constitution Summary

- Never invent facts.
- Always state confidence.
- Cite evidence.
- Admit uncertainty.
- Separate facts from interpretation.
- Challenge consensus.
- Protect capital first.
- Log every decision.
- Recommend no trade when edge is insufficient.

Appendix D - Practical First Implementation

The first build should not be a live trading bot. The first build should be a daily research memo generated by one Macro Agent and one CIO summarizer. After the memo is consistently useful,

add the Technical Agent, then Fundamentals, then Semiconductor, then Risk, then Quant. Paper trading should begin only after the daily memo is stable and the decision log is working.

Bibliography and Influences

Benjamin Graham: The Intelligent Investor; Security Analysis - margin of safety, discipline, and the distinction between investing and speculation.

Warren Buffett / Berkshire Hathaway letters: Business quality, durable competitive advantage, long-term ownership, and management integrity.

Howard Marks: The Most Important Thing and market memos - second-level thinking, cycles, risk awareness, and humility.

Ray Dalio: Principles and macro-cycle thinking - explicit principles, diversification, and systematic decision making.

Harry Markowitz: Portfolio Selection, Journal of Finance, 1952 - portfolio-level risk, covariance, and diversification. [1]

Eugene Fama and Kenneth French: Common Risk Factors in the Returns on Stocks and Bonds, 1993; Five-Factor Asset Pricing Model - factor awareness and performance attribution. [2][3]

Narasimhan Jegadeesh and Sheridan Titman: Returns to Buying Winners and Selling Losers, 1993 - evidence for intermediate-term momentum. [4]

Mark Carhart: On Persistence in Mutual Fund Performance, 1997 - four-factor attribution and momentum in performance analysis. [5]

Daniel Kahneman; Richard Thaler: Behavioral finance, decision errors, overconfidence, loss aversion, and the need for process safeguards.

SEC EDGAR APIs: Official access to company submissions and XBRL company facts. [16]

OpenAI Agents documentation: Agent orchestration, tools, collaboration, structured state, and guardrails. [17]

Alpaca and Interactive Brokers API documentation: Paper trading and broker API capabilities for controlled testing and execution. [18][19]

FINRA / SEC margin guidance: Regulatory constraints around frequent intraday trading and margin requirements; 2026 FINRA changes replace the older PDT framework with intraday margin standards. [20][21]

Source Notes

[1] Markowitz, H. Portfolio Selection. The Journal of Finance, 1952.

[2] Fama, E. and French, K. Common Risk Factors in the Returns on Stocks and Bonds. Journal of Financial Economics, 1993.

[3] Fama, E. and French, K. A Five-Factor Asset Pricing Model. SSRN / Journal of Financial Economics version.

- [4] Jegadeesh, N. and Titman, S. Returns to Buying Winners and Selling Losers: Implications for Stock Market Efficiency. *Journal of Finance*, 1993.
- [5] Carhart, M. On Persistence in Mutual Fund Performance. *Journal of Finance*, 1997.
- [16] SEC EDGAR Application Programming Interfaces documentation.
- [17] OpenAI Agents SDK documentation.
- [18] Alpaca paper trading and trading API documentation.
- [19] Interactive Brokers API documentation.
- [20] FINRA Regulatory Notice 26-10 and FINRA intraday margin explanations, 2026.
- [21] SEC / Investor.gov margin and day-trading investor bulletins.

Disclaimer

This document is educational and operational planning material. It is not personalized investment advice, legal advice, tax advice, or a solicitation to buy or sell securities. Active trading involves risk, including the risk of loss greater than expected. Any live implementation should be reviewed with qualified financial, legal, tax, and compliance professionals.